

Stay safe, friends. Learn to code from home. Use our free 2,000 hour curriculum.

20 JANUARY 2020 / #APACHE KAFKA

How to implement Change Data Capture using Kafka Streams



Luca Florio

PhD, passionate about Distributed Systems. Functional Programming Enthusiast. Krav Maga black belt.



Change Data Capture (CDC) involves observing the changes happening in a database and making them available in a form that can be exploited by other systems.

One of the most interesting use-cases is to make them available as a stream of events. This means you can, for example, catch the events and update a search index as the data are written to the database.

Interesting right? Let's see how to implement a CDC system that can observe the changes made to a NoSQL database (**MongoDB**), stream them through a message broker (**Kafka**), process the messages of the stream (**Kafka Streams**), and update a search index (**Elasticsearch**)!



TL;DR

The full code of the project is available on GitHub in this [repository](#). If you want to skip all my jibber jabber and just run the example, go straight to the

How to run the project section near the end of the article! 😊

Use case & infrastructure

We run a web application that stores photos uploaded by users. People can share their shots, let others download them, create albums, and so on. Users can also provide a description of their photos, as well as Exif metadata and other useful information.

engine. We will focus on this part of our system that is depicted in the following diagram.

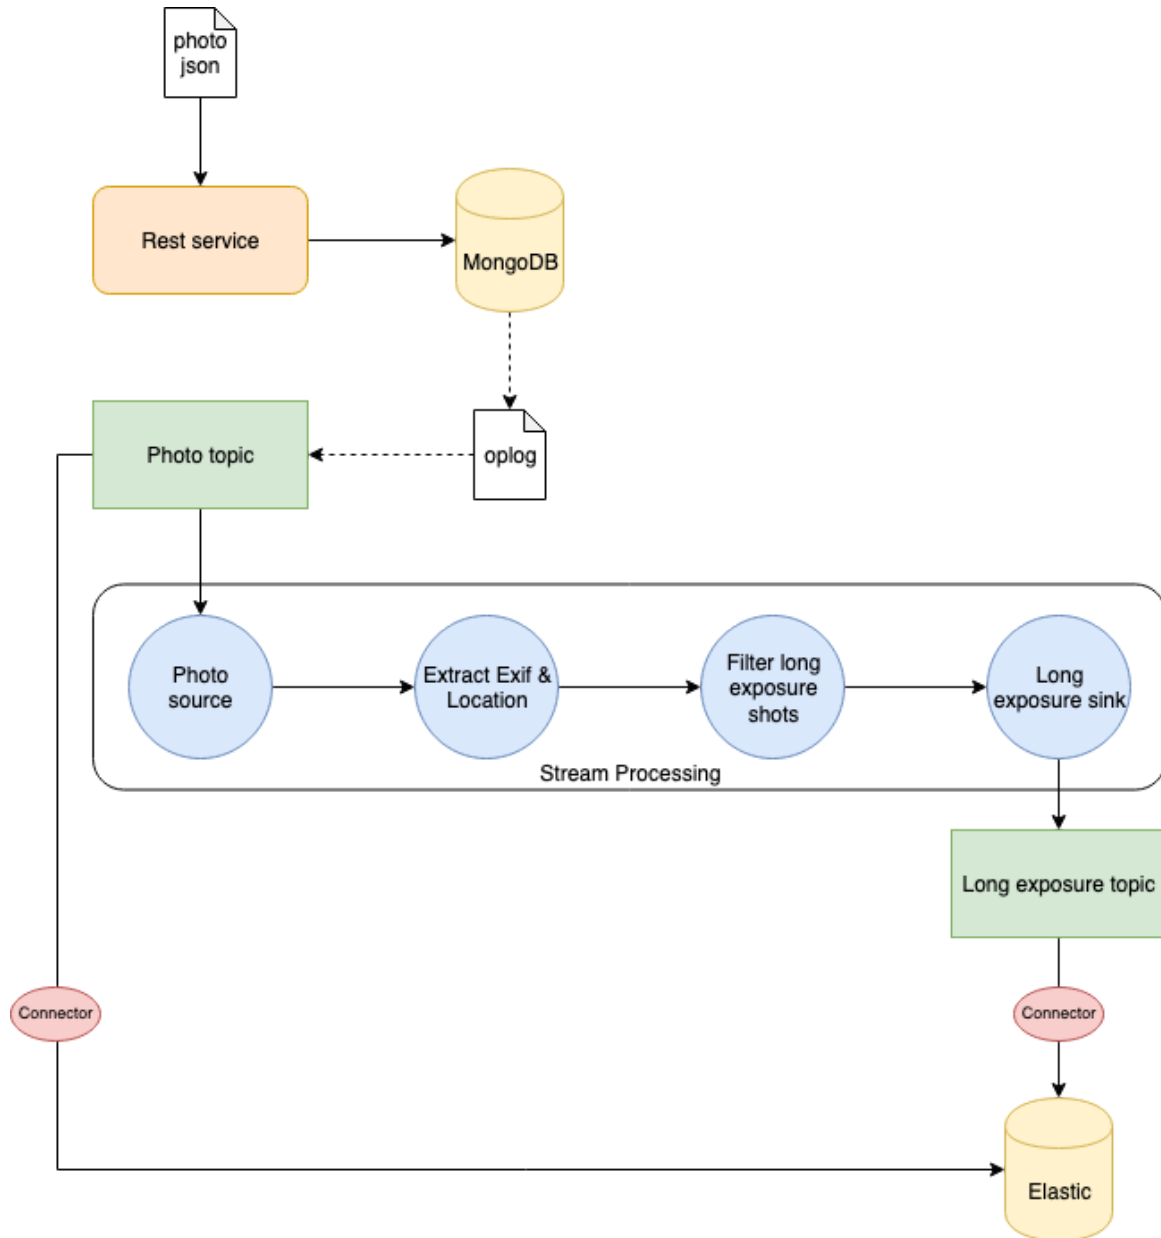


Photo information storage architecture

The information is provided in JSON format. Since I like to post my shots on [Unsplash](#), and the website provides free access to its API, I used their model for the photo JSON document.

Once the JSON is sent through a POST request to our server, we store the document inside a **MongoDB** database. We will also store it in **Elasticsearch** for indexing and quick search.

However, we love **long exposure shots**, and we would like to store in a separate index a subset of information regarding this kind of photo. It can be the exposure time, as well as the location (latitude and longitude) where the photo has been taken. In this way, we can create a map of locations where photographers usually take long exposure photos.

Here comes the interesting part: instead of explicitly calling Elasticsearch in our code once the photo info is stored in MongoDB, we can implement a **CDC** exploiting Kafka and **Kafka Streams**.

We listen to modifications to MongoDB **oplog** using the interface provided by MongoDB itself. When the photo is stored we send it to a photo Kafka topic. Using **Kafka Connect**, an Elasticsearch sink is configured to save everything sent to that topic to a specific index. In this way, we can index all photos stored in MongoDB automatically.

We need to take care of the long exposure photos too. It requires

reason, we use Kafka Streams to create a **processing topology** to:

1. Read from the `photo` topic
2. Extract Exif and location information
3. Filter long exposure photos (exposure time > 1 sec.)
4. Write to a `long-exposure` topic.

Then another Elasticsearch sink will read data from the `long-exposure` topic and write it to a specific index in Elasticsearch.

It is quite simple, but it's enough to have fun with CDC and Kafka Streams! 😊

Server implementation

Let's have a look at what we need to implement: our server exposing the **REST APIs**!

Models and DAO

First things first, we need a model of our data and a **Data Access Object (DAO)** to talk to our MongoDB database.

As I said, the model for the photo JSON information is the one used by Unsplash. Check out the free API [documentation](#) for an example of the JSON we will use.

I created the mapping for the serializaion/deserialization of the photo JSON using [spray-json](#). I'll skip the details about this, if you are curious just look at the [repo](#)!

Let's focus on the model for the long exposure photo.

```
case class LongExposurePhoto(id: String, exposureTime: Float, createdAt:

object LongExposurePhotoJsonProtocol extends DefaultJsonProtocol {
  implicit val longExposurePhotoFormat: RootJsonFormat[LongExposurePhoto]
}
```

This is quite simple: we keep from the photo JSON the information about the `id`, the exposure time (`exposureTime`), when the photo has been created (`createdAt`), and the location where it has been taken. The location comprehends the city, the country, and the position composed of latitude and longitude.

The DAO consists of just the `PhotoDao.scala` class.

```
class PhotoDao(database: MongoDBDatabase, photoCollection: String) {

  val collection: MongoCollection[Document] = database.getCollection(photoCollection)

  def createPhoto(photo: Photo): Future[String] = {
    val doc = Document(photo.toJson.toString())
    doc.put("_id", photo.id)
    collection.insertOne(doc).toFuture()
      .map(_ => photo.id)
  }
}
```

Since I want to keep this example minimal and focused on the CDC implementation, the DAO has just one method to create a new photo

It is straightforward: create a document from the photo JSON, and insert it in mongo using `id` as the one of the photo itself. Then, we can return the `id` of the photo just inserted in a `Future` (the MongoDB API is async).

Kafka Producer

Once the photo is stored inside MongoDB, we have to send it to the `photo` Kafka topic. This means we need a producer to write the message in its topic. The `PhotoProducer.scala` class looks like this.

```
case class PhotoProducer(props: Properties, topic: String) {

  createKafkaTopic(props, topic)
  val photoProducer = new KafkaProducer[String, String](props)

  def sendPhoto(photo: Photo): Future[RecordMetadata] = {
    val record = new ProducerRecord[String, String](topic, photo.id, photo.data)
    photoProducer.send(record)
  }

  def closePhotoProducer(): Unit = photoProducer.close()
}
```

I would say that this is pretty self-explanatory. The most interesting part is probably the `createKafkaTopic` method that is implemented in the `utils` package.

```
def createKafkaTopic(props: Properties, topic: String): Unit = {
  val adminClient = AdminClient.create(props)
  val photoTopic = new NewTopic(topic, 1, 1)
}
```

This method creates the topic in Kafka setting 1 as a partition and replication factor (it is enough for this example). It is not required, but creating the topic in advance lets Kafka balance partitions, select leaders, and so on. This will be useful to get our stream topology ready to process as we start our server.

Event Listener

We have the DAO that writes in MongoDB and the producer that sends the message in Kafka. We need to glue them together in some way so that when the document is stored in MongoDB the message is sent to the `photo` topic. This is the purpose of the `PhotoListener.scala` class.

```
case class PhotoListener(collection: MongoCollection[Document], producer:

    val cursor: ChangeStreamObservable[Document] = collection.watch()

    cursor.subscribe(new Observer[ChangeStreamDocument[Document]] {
      override def onNext(result: ChangeStreamDocument[Document]): Unit = {
        result.getOperationType match {
          case OperationType.INSERT => {
            val photo = result.getFullDocument.toJson().parseJson.convertTo
            producer.sendPhoto(photo).get()
            println(s"Sent photo with Id ${photo.id}")
          }
          case _ => println(s"Operation ${result.getOperationType} not supp
        }
      }
    })
    override def onError(e: Throwable): Unit = println(s"onError: $e")
    override def onComplete(): Unit = println("onComplete")})
}
```


We exploit the Change Streams interface provided by the MongoDB scala library.

Here is how it works: we `watch()` the collection where photos are stored. When there is a new event (`onNext`) we run our logic.

For this example we are interested only in the creation of new documents, so we explicitly check that the operation is of type `OperationType.INSERT`. If the operation is the one we are interested in, we get the document and convert it to a `Photo` object to be sent by our producer.

That's it! With few lines of code we connected the creation of documents in MongoDB to a stream of events in Kafka. 🐘

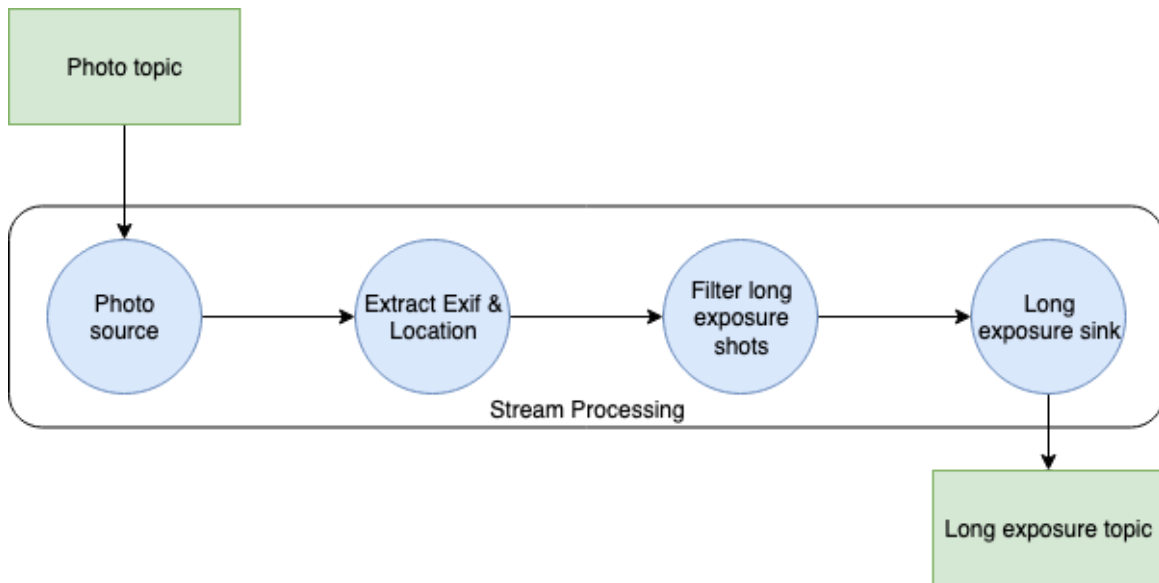
As a side note, be aware that to use the Change Streams interface **we have to setup a MongoDB replica set**. This means we need to run 3 instances of MongoDB and configure them to act as a replica set using the following command in mongo client:

```
rs.initiate({_id : "r0", members: [{ _id : 0, host : "mongo1:27017", pri
```

Here our instances are the containers we will run in the docker-compose file, that is `mongo1`, `mongo2`, and `mongo3`.

Processing Topology

Time to build our processing topology! It will be in charge of the creation of the `long-exposure` index in Elasticsearch. The topology is



Processing topology

and it is implemented in the `LongExposureTopology.scala` object class.

Let's analyse every step of our processing topology.

```
val stringSerde = new StringSerde

val streamsBuilder = new StreamsBuilder()

val photoSource: KStream[String, String] = streamsBuilder.stream(sourceTo
```

The first step is to read from a source topic. We start a stream from the `sourceTopic` (that is `photo` topic) using the `StreamsBuilder()` object. The `stringSerde` object is used to serialise and deserialise the content of the topic as a `String`.

Please notice that at each step of the processing we create a new stream of data with a `KStream` object. When creating the stream, we specify the key and the value produced by the stream. In our topology the key will always be a `String`. In this step the value produced is still a `String`.

```
val covertToPhotoObject: KStream[String, Photo] =  
    photoSource.mapValues((_, jsonString) => {  
        val photo = jsonString.parseJson.convertTo[Photo]  
        println(s"Processing photo ${photo.id}")  
        photo  
    })
```

The next step is to convert the value extracted from the `photo` topic into a proper `Photo` object.

So we start from the `photoSource` stream and work on the values using the `mapValues` function. We simply parse the value as a `JSON` and create the `Photo` object that will be sent in the `covertToPhotoObject` stream.

```
val filterWithLocation: KStream[String, Photo] = covertToPhotoObject.filter
```

info about the location, but we want it in our long exposure object. This step of the topology filters out from the `convertToPhotoObject` stream the photos that have no info about the location, and creates the `filterWithLocation` stream.

```
val filterWithExposureTime: KStream[String, Photo] = filterWithLocation.f
```

Another important fact for our processing is the exposure time of the photo. For this reason, we filter out from the `filterWithLocation` stream the photos without exposure time info, creating the `filterWithExposureTime`.

```
val dataExtractor: KStream[String, LongExposurePhoto] =  
    filterWithExposureTime.mapValues((_, photo) => LongExposurePhoto(pl
```

We now have all we need to create a `LongExposurePhoto` object! That is the result of the `dataExtractor`: it takes the `Photo` coming from the `filterWithExposureTime` stream and produces a new stream containing `LongExposurePhoto`.

```
val longExposureFilter: KStream[String, String] =  
    dataExtractor.filter((_, item) => item.exposureTime > 1.0).mapValue  
    val jsonString = longExposurePhoto.toJson.compactPrint  
    println(s"completed processing: $jsonString")  
    jsonString  
    })
```

We are almost there. We now have to keep the photos with a long exposure time (that we decided is more than 1 sec.). So we create a new `longExposureFilter` stream without the photos that are not long exposure.

This time we also serialise the `LongExposurePhotos` into the corresponding `JSON` string, which will be written to Elasticsearch in the next step.

```
longExposureFilter.to(sinkTopic, Produced.`with`(stringSerde, stringSerde)  
  
streamsBuilder.build()
```

This is the last step of our topology. We write to our `sinkTopic` (that is long-exposure topic) using the string serialiser/deserialiser what is inside the `longExposureFilter` stream.

The last command simply `build`s the topology we just created.

Now that we have our topology, we can use it in our server. The `PhotoStreamProcessor.scala` class is what manages the processing.

```
class PhotoStreamProcessor(kafkaProps: Properties, streamProps: Properties)  
  
  createKafkaTopic(kafkaProps, sinkTopic)  
  val topology: Topology = LongExposureTopology.build(sourceTopic, sinkTopic)  
  val streams: KafkaStreams = new KafkaStreams(topology, streamProps)  
  
  sys.ShutdownHookThread {  
    streams.close(java.time.Duration.ofSeconds(10))  
  }
```

```
def start(): Unit = new Thread {  
  override def run(): Unit = {  
    streams.cleanUp()  
    streams.start()  
    println("Started long exposure processor")  
  }  
}.start()  
}
```

First we create the `sinkTopic`, using the same utility method we saw before. Then we build the stream topology and initialize a `KafkaStreams` object with that topology.

To start the stream processing, we need to create a dedicated `Thread` that will run the streaming while the server is alive. According to the official documentation, it is always a good idea to `cleanUp()` the stream before starting it.

Our `PhotoStreamProcessor` is ready to go! 🍷

REST API

The server exposes REST APIs to send it the photo information to store. We make use of [Akka HTTP](#) for the API implementation.

```
trait AppRoutes extends SprayJsonSupport {  
  
  implicit def system: ActorSystem  
  implicit def photoDao: PhotoDao  
  implicit lazy val timeout = Timeout(5.seconds)  
  
  lazy val healthRoute: Route = pathPrefix("health") {  
    concat(  
      pathEnd {
```



```

        complete(StatusCodes.OK)
      }
    )
  }
)
}

lazy val crudRoute: Route = pathPrefix("photo") {
  concat(
    pathEnd {
      concat(
        post {
          entity(as[Photo]) { photo =>
            val photoCreated: Future[String] =
              photoDao.createPhoto(photo)
            onSuccess(photoCreated) { id =>
              complete((StatusCodes.Created, id))
            }
          }
        }
      )
    }
  )
}

```

To keep the example minimal, we have only two routes:

- GET /health - to check if the server is up & running
- POST /photo - to send to the system the JSON of the photo information we want to store. This endpoint uses the DAO to store the document in MongoDB and returns a 201 with the id of the stored photo if the operation succeeded.

This is by no means a complete set of APIs, but it is enough to run our example. 😊

Server main class

OK, we implemented all the components of our server, so it's time to wrap everything up. This is our `Server.scala` object class.

```
implicit val system: ActorSystem = ActorSystem("kafka-stream-playground")
implicit val materializer: ActorMaterializer = ActorMaterializer()
```

First a couple of **Akka** utility values. Since we use Akka HTTP to run our server and REST API, these implicit values are required.

```
val config: Config = ConfigFactory.load()
val address = config.getString("http.ip")
val port = config.getInt("http.port")

val mongoUri = config.getString("mongo.uri")
val mongoDb = config.getString("mongo.db")
val mongoUser = config.getString("mongo.user")
val mongoPwd = config.getString("mongo.pwd")
val photoCollection = config.getString("mongo.photo_collection")

val kafkaHosts = config.getString("kafka.hosts").split(',').toList
val photoTopic = config.getString("kafka.photo_topic")
val longExposureTopic = config.getString("kafka.long_exposure_topic")
```

Then we read all the configuration properties. We will come back to the configuration file in a moment.

```
val kafkaProps = new Properties()
kafkaProps.put("bootstrap.servers", kafkaHosts.mkString(","))
```

```
val streamProps = new Properties()
streamProps.put(StreamsConfig.APPLICATION_ID_CONFIG, "long-exp-proc-app")
streamProps.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, kafkaHosts.mkStri

val photoProducer = PhotoProducer(kafkaProps, photoTopic)
val photoStreamProcessor = new PhotoStreamProcessor(kafkaProps, streamProc
photoStreamProcessor.start()
```

We have to configure both our Kafka producer and the stream processor. We also start the stream processor, so the server will be ready to process the documents sent to it.

```
val client = MongoClient(s"$mongodb://$mongoUri/$mongoUser")
val db = client.getDatabase(mongoDb)
val photoDao: PhotoDao = new PhotoDao(db, photoCollection)
val photoListener = PhotoListener(photoDao.collection, photoProducer)
```

Also MongoDB needs to be configured. We setup the connection and initialize the DAO as well as the listener.

```
lazy val routes: Route = healthRoute ~ crudRoute

Http().bindAndHandle(routes, address, port)
Await.result(system.whenTerminated, Duration.Inf)
```

Everything has been initialized. We create the REST routes for the communication to the server, bind them to the handlers, and finally

Server configuration

This is the configuration file used to setup the server:

```
http {
  ip = "127.0.0.1"
  ip = ${?SERVER_IP}

  port = 8000
  port = ${?SERVER_PORT}
}
mongo {
  uri = "127.0.0.1:27017"
  uri = ${?MONGO_URI}
  db = "kafka-stream-playground"
  user = "admin"
  pwd = "admin"
  photo_collection = "photo"
}
kafka {
  hosts = "127.0.0.1:9092"
  hosts = ${?KAFKA_HOSTS}
  photo_topic = "photo"
  long_exposure_topic = "long-exposure"
}
```

I think that this one does not require much explanation, right? 😊

Connectors configuration

The server we implemented writes in two Kafka topics: `photo` and `long-exposure`. But how are messages written in Elasticsearch as documents? Using **Kafka Connect**!

to write every message going through that topic in Elasticsearch.

First we need Kafka Connect. We can use the container provided by Confluence in the docker-compose file:

```
connect:
  image: confluentinc/cp-kafka-connect
  ports:
    - 8083:8083
  networks:
    - kafka_stream_playground
  depends_on:
    - zookeeper
    - kafka
  volumes:
    - $PWD/connect-plugins:/connect-plugins
  environment:
    CONNECT_BOOTSTRAP_SERVERS: kafka:9092
    CONNECT_REST_ADVERTISED_HOST_NAME: connect
    CONNECT_REST_PORT: 8083
    CONNECT_GROUP_ID: compose-connect-group
    CONNECT_CONFIG_STORAGE_TOPIC: docker-connect-configs
    CONNECT_CONFIG_STORAGE_REPLICATION_FACTOR: 1
    CONNECT_OFFSET_FLUSH_INTERVAL_MS: 10000
    CONNECT_OFFSET_STORAGE_TOPIC: docker-connect-offsets
    CONNECT_OFFSET_STORAGE_REPLICATION_FACTOR: 1
    CONNECT_STATUS_STORAGE_TOPIC: docker-connect-status
    CONNECT_STATUS_STORAGE_REPLICATION_FACTOR: 1
    CONNECT_KEY_CONVERTER: "org.apache.kafka.connect.storage.StringConverter"
    CONNECT_VALUE_CONVERTER: "org.apache.kafka.connect.json.JsonConverter"
    CONNECT_VALUE_CONVERTER_SCHEMAS_ENABLE: "false"
    CONNECT_INTERNAL_KEY_CONVERTER: "org.apache.kafka.connect.json.JsonConverter"
    CONNECT_INTERNAL_VALUE_CONVERTER: "org.apache.kafka.connect.json.JsonConverter"
    CONNECT_ZOOKEEPER_CONNECT: zookeeper:2181
    CONNECT_PLUGIN_PATH: /connect-plugins
    CONNECT_LOG4J_ROOT_LOGLEVEL: INFO
```

I want to focus on some of the configuration values.

Therefore, we need to expose the port 8083 and will use the endpoint to configure the connectors (`CONNECT_REST_PORT`).

We also need to map a volume to the `/connect-plugins` path, where we will place the Elasticsearch Sink Connector to write to Elasticsearch. This is reflected also in the `CONNECT_PLUGIN_PATH` .

The `connect` container should know how to find the Kafka servers, so we set `CONNECT_BOOTSTRAP_SERVERS` as `kafka:9092` .

Once Kafka Connect is ready, we can send the configurations of our connectors to the `http://localhost:8083/connectors` endpoint. We need 2 connectors, one for the `photo` topic and one for the `long-exposure` topic. We can send the configuration as a `JSON` with a `POST` request.

```
{
  "name": "photo-connector",
  "config": {
    "connector.class": "io.confluent.connect.elasticsearch.ElasticsearchSinkConnector",
    "tasks.max": "1",
    "topics": "photo",
    "key.converter": "org.apache.kafka.connect.storage.StringConverter",
    "value.converter": "org.apache.kafka.connect.json.JsonConverter",
    "value.converter.schemas.enable": "false",
    "schema.ignore": "true",
    "connection.url": "http://elastic:9200",
    "type.name": "kafka-connect",
    "behavior.on.malformed.documents": "warn",
    "name": "photo-connector"
  }
}
```


as the `connector.class` , as well as the `topics` that we want to sink - in this case `photo` .

We don't want to use a schema for the `value.converter` , so we can disable it (`value.converter.schemas.enable`) and tell the connector to ignore the schema (`schema.ignore`).

The connector for the `long-exposure` topic is exactly like this one. The only difference is the `name` and of course the `topics` .

How to run the project

We have all we need to test the CDC! How can we do it? It's quite easy: simply run the `setup.sh` script in the root folder of the repo!

What will the script do?

1. Run the `docker-compose` file with all the services.
2. Configure MongoDB replica set. This is required to enable the **Change Stream interface** to capture data changes. More info about this [here](#).
3. Configure the Kafka connectors.
4. Connect to the logs of the server.

The docker-compose will run the following services:

- Our Server
- 3 instances of MongoDB (required for the replica set)

- Kafka (single node)
- Kafka connect
- Zookeeper (required by Kafka)
- Elasticsearch
- Kibana

There are a lot of containers to run, so make sure you have enough resources to run everything properly. If you want, remove Mongoku and Kibana from the compose-file, since they are used just for a quick look inside the DBs.

Once everything is up and running, you just have to send data to the server.

I collected some JSON documents of photos from Unplash that you can use to test the system in the `photos.txt` file.

There are a total of 10 documents, with 5 of them containing info about long exposure photos. Send them to the server running the `send-photos.sh` script in the root of the repo. Check that everything is stored in MongoDB connecting to Mongoku at `http://localhost:3100`. Then connect to Kibana at `http://localhost:5601` and you will find two indexes in Elasticsearch: `photo`, containing the JSON of all the photos stored in MongoDB, and `long-exposure`, containing just the info of the long exposure photos.

Amazing, right? 😊

Conclusion

Starting from the design of the use-case, we built our system that connected a MongoDB database to Elasticsearch using CDC.

Kafka Streams is the enabler, allowing us to convert database events to a stream that we can process.

Do you need to see the whole project? Just checkout the [repository](#) on GitHub! 😊

That's it, enjoy! 🚀

If you read this far, tweet to the author to show them you care.

[Tweet a thanks](#)

Learn to code for free. freeCodeCamp's open source curriculum has helped more than 40,000 people get jobs as developers.

[Get started](#)

freeCodeCamp is a donor-supported tax-exempt 501(c)(3) nonprofit organization (United States Federal Tax Identification Number: 82-0779546)

Our mission: to help people learn to code for free. We accomplish this by creating thousands of videos, articles, and interactive coding lessons - all freely available to the public. We also have thousands of freeCodeCamp study groups around the world.

Donations to freeCodeCamp go toward our education initiatives, and help pay for servers, services, and staff

[you can make a tax-deductible donation here.](#)

Our Nonprofit

[About](#)[Alumni Network](#)[Open Source](#)[Shop](#)[Support](#)[Sponsors](#)[Academic Honesty](#)[Code of Conduct](#)[Privacy Policy](#)[Terms of Service](#)[Copyright Policy](#)

Trending Guides

[2019 Web Developer Roadmap](#)[Python Tutorial](#)[CSS Flexbox Guide](#)[JavaScript Tutorial](#)[Python Example](#)[HTML Tutorial](#)[Linux Command Line Guide](#)[JavaScript Example](#)[Git Tutorial](#)[React Tutorial](#)[Java Tutorial](#)[Linux Tutorial](#)[CSS Tutorial](#)[jQuery Example](#)[SQL Tutorial](#)[CSS Example](#)[React Example](#)[Angular Tutorial](#)[Bootstrap Example](#)[How to Set Up SSH Keys](#)[WordPress Tutorial](#)[PHP Example](#)

